

COLLECTIVE VARIABLES MODULE

reference manual for LAMMPS

June 4, 2013



Giacomo Fiorin, Jérôme Hénin, Axel Kohlmeyer

Contents

1	Introduction	3
2	General parameters and input/output files	3
2.1	LAMMPS keywords	3
2.2	Configuration file for the collective variables module	5
2.3	Input state file (optional)	7
2.4	Output files	7
3	Defining collective variables and their properties	8
3.1	General options for a collective variable	8
3.2	Artificial boundary potentials (walls)	9
3.3	Trajectory output	10
3.4	Extended Lagrangian.	11
3.5	Statistical analysis of collective variables	12
4	Selecting atoms for colvars: defining atom groups	13
4.1	Selection keywords	13
4.2	Moving frame of reference.	15
4.3	Treatment of periodic boundary conditions.	16
4.4	Computational cost of colvars based on group size.	16
5	Collective variable components (basis functions)	17
5.1	List of available colvar components	18
5.1.1	distance: center-of-mass distance between two groups.	18
5.1.2	distanceZ: projection of a distance vector on an axis.	18
5.1.3	distanceXY: modulus of the projection of a distance vector on a plane.	19
5.1.4	distanceVec: distance vector between two groups.	20
5.1.5	distanceDir: distance unit vector between two groups.	20
5.1.6	distanceInv: mean distance between two groups of atoms.	20
5.1.7	angle: angle between three groups.	20
5.1.8	dihedral: torsional angle between four groups.	20
5.1.9	coordNum: coordination number between two groups.	21
5.1.10	selfCoordNum: coordination number between atoms within a group.	22
5.1.11	rmsd: root mean square displacement (RMSD) from reference positions.	22
5.1.12	Advanced usage of the rmsd component.	23
5.1.13	eigenvector: projection of the atomic coordinates on a vector.	23
5.1.14	gyration: radius of gyration of a group of atoms.	24
5.1.15	inertia: total moment of inertia of a group of atoms.	24
5.1.16	inertiaZ: total moment of inertia of a group of atoms around a chosen axis.	24
5.1.17	orientation: orientation from reference coordinates.	25
5.1.18	orientationAngle: angle of rotation from reference coordinates.	25
5.1.19	spinAngle: angle of rotation around a given axis.	25
5.1.20	tilt: cosine of the rotation orthogonal to a given axis.	26
5.2	Advanced usage and special considerations	26
5.2.1	Periodic components.	26
5.2.2	Non-scalar components.	27

5.2.3	Calculating system forces.	27
5.3	Linear and polynomial combinations of components	28
6	Biasing and analysis methods	28
6.1	Adaptive Biasing Force	29
6.1.1	ABF requirements on collective variables	30
6.1.2	Output files	32
6.1.3	Post-processing: reconstructing a multidimensional free energy surface	32
6.2	Metadynamics	33
6.2.1	Output files	34
6.2.2	Performance tuning	35
6.2.3	Well-tempered metadynamics	36
6.2.4	Multiple-replicas metadynamics	36
6.2.5	Compatibility and post-processing	37
6.3	Harmonic restraints	38
6.3.1	Moving restraints: steered molecular dynamics	39
6.3.2	Moving restraints: umbrella sampling	40
6.3.3	Changing force constant	40
6.4	Multidimensional histograms	41

1 Introduction

In today’s molecular dynamics simulations, it is often useful to reduce the great number of degrees of freedom of a into a few parameters which can be either analyzed individually, or manipulated in order to alter the dynamics in a controlled manner. These have been called ‘order parameters’, ‘collective variables’, ‘(surrogate) reaction coordinates’, and many other terms. Here we use primarily the term ‘collective variable’ (shortened to *colvar*), which indicates any differentiable function of atomic Cartesian coordinates, x_i , with i between 1 and N , the total number of atoms:

$$\xi(t) = \xi(x_i(t), x_j(t), x_k(t), \dots), \quad 1 \leq i, j, k \dots \leq N \quad (1)$$

This manual documents the collective variables module (colvars), a portable software that interfaces multiple MD simulation simulation programs, with a focus on flexibility, robustness and high performance. The module is designed to perform multiple tasks concurrently, the most common of which are:

- apply restraints or biasing potentials to multiple colvars, tailored on the system by choosing from a wide set of basis functions, without limitations on their number or on the number of atoms involved;
- calculate potentials of mean force (PMFs) along any set of colvars, using different enhanced sampling methods, such as Adaptive Biasing Force (ABF), metadynamics, steered MD and umbrella sampling; variants of these methods that make use of an ensemble of replicas are supported as well;
- calculate statistical properties of the colvars, such as running averages and standard deviations, correlation functions of pairs of colvars, and multidimensional histograms, without the need to save very large trajectory files.

To briefly illustrate the flexibility of the colvars module, Figure 1 shows an example of a non-trivial configuration (the corresponding input can be found in 2.2).

2 General parameters and input/output files

To enable a collective variable calculation, one or more parameters are added to the LAMMPS configuration file. In 2.1, we document the syntax of such parameters. One of these is the name of a configuration file for the collective variables module itself, which is described in 2.2, and in the following sections.

2.1 LAMMPS keywords

To enable a collective variables-based calculation, the following line must be added to the LAMMPS configuration file:

```
fix ID group-ID colvars configfile keyword values ...
```

where *ID* and *group-ID* are general options of the `fix` command, *configfile* is the name of the configuration file for the collective variables module, and *keyword* is one of the following optional keywords:

- `input` <Name or prefix of the input state file>
Context: Keyword of the `fix colvars` command
Acceptable values: string
Description: If a value is provided, it is interpreted as either the name of the input state file, or as the prefix of the file named *input.colvars.state*. This allows to continue a previous collective variables-based calculation (see 2.3).



Figure 1: Graphical representation of a collective variables configuration. The colvar called “*d*” is defined as the difference between two distances: the first distance (d_1) is taken between the center of mass of atoms 1 and 2 and that of atoms 3 to 5, the second (d_2) between atom 7 and the center of mass of atoms 8 to 10. The difference $d = d_1 - d_2$ is obtained by multiplying the two by a coefficient $C = +1$ or $C = -1$, respectively. The colvar called “*c*” is the coordination number calculated between atoms 1 to 10 and atoms 11 to 20. A harmonic restraint is applied to both d and c : to allow using the same force constant K , both d and c are scaled by their respective fluctuation widths w_d and w_c . A third colvar “*alpha*” is defined as the α -helical content of residues 1 to 10. The values of “*c*” and “*alpha*” are also recorded throughout the simulation as a joint 2-dimensional histogram. **Note:** currently, the α -helical content colvar is unavailable in LAMMPS, as it requires a name-based topology; future releases will overcome this limitation.

- output \langle Prefix of the output state file \rangle
Context: Keyword of the `fix colvars` command
Acceptable values: string
Default value: “out”
Description: If a value is provided, it is interpreted as the prefix to all output files that will be written by the collective variables module (see 2.4).
- unwrap \langle Whether to unwrap coordinates passed to the colvars module \rangle
Context: keyword of the `fix colvars` command
Acceptable values: “yes” or “no”
Default value: “yes”
Description: This keyword controls whether wrapped or unwrapped coordinates are passed to the colvars module for calculation of the collective variables and of the resulting forces. The default is to use the image flags to reconstruct the absolute atom positions: under this convention, centers of mass and centers of geometry are calculated as a weighted vector sum (see 4.3). Setting this to no will use the current local coordinates that are wrapped back into the simulation cell at each re-neighboring instead.

- **seed** \langle Seed for the random number generator \rangle
Context: Keyword of the `fix colvars` command
Acceptable values: positive integer
Default value: 1966
Description: If defined, the value of this keyword is provided as seed to the random number generator. This is only meaningful when the `extendedLangevinDamping` keyword is used (see 3.4).
- **tstat** \langle Thermostating fix \rangle
Context: Keyword of the `fix colvars` command
Acceptable values: string
Default value: NULL
Description: This keyword provides the label of a thermostating fix that thermostats all atoms in the `fix colvars` group. This will be used to provide the `colvars` module with the current thermostat target temperature.

2.2 Configuration file for the collective variables module

The syntax of the `colvars` configuration file is “keyword value”, where the keyword and its value are separated by any white space. The following rules apply:

- keywords are case-insensitive (`upperBoundary` is the same as `upperboundary` and `UPPERBOUNDARY`); their string values are however case-sensitive (e.g. file names);
- a long value or a list of multiple values can be distributed across multiple lines by using curly braces, “{” and “}”: the opening brace “{” must occur on the same line as the keyword, following a space character or other white space; the closing brace “}” can be at any position after that;
- many keywords are nested, and are only meaningful within a specific context: for every keyword documented in the following, the “parent” keyword that defines such context is also indicated;
- if a keyword requiring a boolean value (`yes|on|true` or `no|off|false`) is provided without an explicit value, it defaults to ‘yes|on|true’; for example, ‘`outputAppliedForce`’ may be used as shorthand for ‘`outputAppliedForce on`’;
- the hash character `#` indicates a comment: all text in the same line following this character will be ignored.

All parameters defining the `colvars` and their biasing or analysis algorithms are read from the file specified by `colvarsConfig`. The following keywords are available in the global context of this configuration file, i.e. not nested inside other keywords:

- **colvarsTrajFrequency** \langle Colvar value trajectory frequency \rangle
Context: global
Acceptable values: positive integer
Default value: 100
Description: The values of each colvar (and of other related quantities, if requested) are written to the file `output.colvars.traj` every these many steps throughout the simulation. If the value is 0, such trajectory file is not written. For optimization the output is buffered, and synchronized with the disk only when the restart file is being written.

- `colvarsTrajAppend` `<Append to trajectory file?>`
Context: global
Acceptable values: boolean
Default value: off
Description: If this flag is enabled, and a file with the same name as the trajectory file is already present, new data is appended to that file. Otherwise, a new file is created with the same name that overwrites the previous file.
- `colvarsRestartFrequency` `<Colvar module restart frequency>`
Context: global
Acceptable values: positive integer
Default value: `restartFreq`
Description: Allows to choose a different restart frequency for the collective variables module. Redefining it may be useful to trace the time evolution of those few properties which are not written to the trajectory file for reasons of disk space.
- `indexFile` `<Index file for atom selection (GROMACS “ndx” format)>`
Context: global
Acceptable values: UNIX filename
Description: This option reads an index file (usually with a `.ndx` extension) as produced by the `make_ndx` tool of GROMACS. In LAMMPS, the `group2ndx` command can be used to generate such file from existing groups. The names of index groups contained in this file can then be used to define atom groups with the `indexGroup` keyword. Other supported methods to select atoms are described in 4.
- `analysis` `<Turn on run-time statistical analysis >`
Context: global
Acceptable values: boolean
Default value: off
Description: If this flag is enabled, each colvar is instructed to perform whatever run-time statistical analysis it is configured to, such as correlation functions, or running averages and standard deviations. See section 3.5 for details.

The example below defines the same configuration shown in Fig. 1. The options within the colvar blocks are described in 3 and 5, the ones within the harmonic and histogram blocks in 6. **Note:** *except colvar, none of the keywords shown is mandatory.*

```
colvar {
  # difference of two distances
  name d
  width 0.2 # 0.2 of estimated fluctuation width
  distance {
    componentCoeff 1.0
    group1 { atomNumbers 1 2 }
    group2 { atomNumbers 3 4 5 }
  }
  distance {
    componentCoeff -1.0
    group1 { atomNumbers 7 }
    group2 { atomNumbers 8 9 10 }
```

```

    }
}

colvar {
  name c
  coordNum {
    cutoff 6.0
    group1 { atomNumbersRange 1-10 }
    group2 { atomNumbersRange 11-20 }
  }
}

harmonic {
  colvars d c
  centers 3.0 4.0
  forceConstant 5.0
}

histogram {
  colvars c
}

```

Section 3 explains how to define a colvar and its behavior, regardless of its specific functional form. To define colvars that are appropriate to a specific physical system, Section 4 documents how to select atoms, and section 5 lists all of the available functional forms, which we call “colvar components”. Finally, section 6 lists the available methods and algorithms to perform biased simulations and multidimensional analysis of colvars.

2.3 Input state file (optional)

Aside from the configuration file, an optional input state file may be provided to continue a previous simulation. The name of this file is provided as the argument to the input keyword of the `fix ID group-ID colvars` command.

2.4 Output files

In addition to the output files written by the MD program, the following three output files are written:

- a *state file*, named `output.colvars.state`; this file is in ASCII format.
- if the parameter `colvarsRestartFrequency` is larger than zero, a *restart file* named `restart.colvars.state` is written, which is equivalent to the state file;
- if the parameter `colvarsTrajFrequency` is greater than 0 (default: 100), a *trajectory file* is written during the simulation: its name is `output.colvars.traj`; unlike the state file, it is not needed to restart a simulation, but can be used later for post-processing and analysis.

Other output files may be written by specific methods applied to the colvars (e.g. by the ABF method, see 6.1, or the metadynamics method, see 6.2). Like the colvar trajectory file, they are needed only for analyzing, not continuing a simulation. All such files’ names also begin with the prefix *output*.

3 Defining collective variables and their properties

In the configuration file each colvar is defined by the keyword `colvar`, followed by its configuration options within curly braces: `colvar { ... }`. One of these options is the name of a colvar component: for example, including `rmsd { ... }` defines the colvar as a RMSD function. *In most applications, only one component is used, and the component is equal to the colvar.*

The full list of colvar components can be found in Section 5, with the syntax to select atoms in Section 4. The following section lists several options to control the behavior of a single colvar, regardless of its type.

3.1 General options for a collective variable

The following options are not required by default; however, the first four are very frequently used:

- `name` `<Name of this colvar>`
Context: `colvar`
Acceptable values: `string`
Default value: “colvar” + numeric id
Description: The name is an unique case-sensitive string which allows the colvar module to identify this colvar unambiguously; it is also used in the trajectory file to label to the columns corresponding to this colvar.
- `width` `<Expected fluctuations amplitude, and resolution for grid-based methods>`
Context: `colvar`
Acceptable values: `positive decimal`
Default value: 1.0
Description: This number is a user-provided estimate of the fluctuation amplitude for the colvar. For example, it is recommended to set this number smaller than or equal to the standard deviation of the colvar during a very short simulation run. Biasing algorithms use this parameter for different purposes: harmonic restraints (6.3) use it to set the physical unit of the force constant, the histogram (6.4) and ABF biases (6.1) interpret it as the grid spacing in the direction of this variable, and metadynamics (6.2) uses it to set the width of newly added hills. This number is expressed in the same physical unit as the colvar value.
- `lowerBoundary` `<Lower boundary of the colvar>`
Context: `colvar`
Acceptable values: `decimal`
Description: Defines the lowest end of the interval of “relevant” values for the colvar. This number can be either a true physical boundary, or a user-defined number. Together with `upperBoundary` and `width`, it is used to define a grid of values along the colvar (not available for colvars based on `distanceDir`, `distanceVec`, and `orientation`). This option does not affect dynamics: to confine a colvar within a certain interval, the options `lowerWall` and `lowerWallConstant` should be used.
- `upperBoundary` `<Upper boundary of the colvar>`
Context: `colvar`
Acceptable values: `decimal`
Description: Similarly to `lowerBoundary`, defines the highest possible or allowed value.
- `hardLowerBoundary` `<Whether the lower boundary is the physical lower limit>`
Context: `colvar`
Acceptable values: `boolean`

Default value: off

Description: This option does not affect simulation results, but enables some internal optimizations. Depending on its mathematical definition, a colvar may have “natural” boundaries: for example, a distance colvar has a “natural” lower boundary at 0 Å. Setting this option instructs the colvars module that the user-defined lower boundary is “natural”. See Section 5 for the physical ranges of values of each component.

- `hardUpperBoundary` 〈Whether the upper boundary is the physical upper limit of the colvar’s values〉

Context: colvar

Acceptable values: boolean

Default value: off

Description: Analogous to `hardLowerBoundary`.

- `expandBoundaries` 〈Allow to expand the two boundaries if needed〉

Context: colvar

Acceptable values: boolean

Default value: off

Description: If defined, biasing and analysis methods may keep their own copies of `lowerBoundary` and `upperBoundary`, and expand them to accommodate values that do not fit in the initial range. Currently, this option is used by the metadynamics bias (6.2) to keep all of its hills fully within the grid. This option cannot be used when the initial boundaries already span the full period of a periodic colvar.

3.2 Artificial boundary potentials (walls)

The following options are useful to define restraints (confining potentials) for this colvar. To apply moving restraints, or restraints to more than one colvar simultaneously, a more convenient option is to use the harmonic bias (6.3).

- `lowerWallConstant` 〈Lower wall force constant (kcal/mol)〉

Context: colvar

Acceptable values: positive decimal

Description: Defines the force constant for a confining restraint on the colvar, in the form of a “half-harmonic” potential. The potential starts at `lowerWall` if it is defined, or `lowerBoundary` otherwise. The energy unit of the constant is kcal/mol, while the spatial unit is that of the colvar.

- `lowerWall` 〈Position of the lower wall〉

Context: colvar

Acceptable values: decimal

Default value: `lowerBoundary`

Description: Defines the value below which a confining restraint on the colvar is applied, in the form of a “half-harmonic” potential. Allows to use a different position of the wall than `lowerBoundary`.

- `upperWallConstant` 〈Upper wall force constant (kcal/mol)〉

Context: colvar

Acceptable values: positive decimal

Description: Analogous to `lowerWallConstant`.

- upperWall <Position of the upper wall>
Context: colvar
Acceptable values: decimal
Default value: upperBoundary
Description: Analogous to lowerWall.

3.3 Trajectory output

- outputValue <Output a trajectory for this colvar>
Context: colvar
Acceptable values: boolean
Default value: on
Description: If colvarsTrajFrequency is non-zero, the value of this colvar is written to the trajectory file every colvarsTrajFrequency steps in the column labeled “<name>”.
- outputVelocity <Output a velocity trajectory for this colvar>
Context: colvar
Acceptable values: boolean
Default value: off
Description: If colvarsTrajFrequency is defined, the finite-difference calculated velocity of this colvar are written to the trajectory file under the label “v_<name>”.
- outputEnergy <Output an energy trajectory for this colvar>
Context: colvar
Acceptable values: boolean
Default value: off
Description: This option applies only to extended Lagrangian colvars. If colvarsTrajFrequency is defined, the kinetic energy of the extended degree and freedom and the potential energy of the restraining spring are written to the trajectory file under the labels “Ek_<name>” and “Ep_<name>”.
- outputSystemForce <Output a system force trajectory for this colvar>
Context: colvar
Acceptable values: boolean
Default value: off
Description: If colvarsTrajFrequency is defined, the total system force on this colvar (i.e. the projection of all interatomic forces except constraint forces on this colvar — see equation (15) in section 6.1) are written to the trajectory file under the label “fs_<name>”. **Note:** not all components support this option. The physical unit for this force is kcal/mol divided by the colvar unit.
- outputAppliedForce <Output an applied force trajectory for this colvar>
Context: colvar
Acceptable values: boolean
Default value: off
Description: If colvarsTrajFrequency is defined, the total force applied on this colvar by biases within the colvar module are written to the trajectory under the label “fa_<name>”. The physical unit for this force is kcal/mol divided by the colvar unit.

3.4 Extended Lagrangian.

The following options enable extended-system dynamics, where a colvar is coupled to an additional degree of freedom (fictitious particle) by a harmonic spring.

- **extendedLagrangian** 〈Add extended degree of freedom〉
Context: colvar
Acceptable values: boolean
Default value: off
Description: Adds a fictitious particle to be coupled to the colvar by a harmonic spring. The fictitious mass and the force constant of the coupling potential are derived from the parameters `extendedTimeConstant` and `extendedFluctuation`, described below. Biasing forces on the colvar are applied to this fictitious particle, rather than to the atoms directly. This implements the extended Lagrangian formalism used in some metadynamics simulations [1].
- **extendedFluctuation** 〈Standard deviation between the colvar and the fictitious particle (colvar unit)〉
Context: colvar
Acceptable values: positive decimal
Default value: width
Description: Defines the spring stiffness for the `extendedLagrangian` mode, by setting the typical deviation between the colvar and the extended degree of freedom due to thermal fluctuation. The spring force constant is calculated internally as $k_B T / \sigma^2$, where σ is the value of `extendedFluctuation`.
- **extendedTimeConstant** 〈Oscillation period of the fictitious particle (fs)〉
Context: colvar
Acceptable values: positive decimal
Default value: 200
Description: Defines the inertial mass of the fictitious particle, by setting the oscillation period of the harmonic oscillator formed by the fictitious particle and the spring. The period should be much larger than the MD time step to ensure accurate integration of the extended particle's equation of motion. The fictitious mass is calculated internally as $k_B T (\tau / 2\pi\sigma)^2$, where τ is the period and σ is the typical fluctuation (see above).
- **extendedTemp** 〈Temperature for the extended degree of freedom (K)〉
Context: colvar
Acceptable values: positive decimal
Default value: thermostat temperature
Description: Temperature used for calculating the coupling force constant of the extended coordinate (see `extendedFluctuation`) and, if needed, as a target temperature for extended Langevin dynamics (see `extendedLangevinDamping`). This should normally be left at its default value.
- **extendedLangevinDamping** 〈Damping factor for extended Langevin dynamics (ps^{-1})〉
Context: colvar
Acceptable values: positive decimal
Default value: 1.0
Description: If this is non-zero, the extended degree of freedom undergoes Langevin dynamics at temperature `extendedTemp`. The friction force is minus `extendedLangevinDamping` times the velocity. This is useful because the extended dynamics coordinate may heat up in the transient non-equilibrium regime of ABF. Use moderate damping values, to limit viscous friction (potentially slow-

ing down diffusive sampling) and stochastic noise (increasing the variance of statistical measurements). In doubt, use the default value.

3.5 Statistical analysis of collective variables

When the global keyword `analysis` is defined in the configuration file, run-time calculations of statistical properties for individual colvars can be performed. At the moment, several types of time correlation functions, running averages and running standard deviations are available.

- `corrFunc` <Calculate a time correlation function?>
Context: colvar
Acceptable values: boolean
Default value: off
Description: Whether or not a time correlation function should be calculated for this colvar.
- `corrFuncWithColvar` <Colvar name for the correlation function>
Context: colvar
Acceptable values: string
Description: By default, the auto-correlation function (ACF) of this colvar, ξ_i , is calculated. When this option is specified, the correlation function is calculated instead with another colvar, ξ_j , which must be of the same type (scalar, vector, or quaternion) as ξ_i .
- `corrFuncType` <Type of the correlation function>
Context: colvar
Acceptable values: velocity, coordinate or coordinate_p2
Default value: velocity
Description: With coordinate or velocity, the correlation function $C_{i,j}(t) = \langle \Pi(\xi_i(t_0), \xi_j(t_0 + t)) \rangle$ is calculated between the variables ξ_i and ξ_j , or their velocities. $\Pi(\xi_i, \xi_j)$ is the scalar product when calculated between scalar or vector values, whereas for quaternions it is the cosine between the two corresponding rotation axes. With coordinate_p2, the second order Legendre polynomial, $(3\cos(\theta)^2 - 1)/2$, is used instead of the cosine.
- `corrFuncNormalize` <Normalize the time correlation function?>
Context: colvar
Acceptable values: boolean
Default value: on
Description: If enabled, the value of the correlation function at $t = 0$ is normalized to 1; otherwise, it equals to $\langle O(\xi_i, \xi_j) \rangle$.
- `corrFuncLength` <Length of the time correlation function>
Context: colvar
Acceptable values: positive integer
Default value: 1000
Description: Length (in number of points) of the time correlation function.
- `corrFuncStride` <Stride of the time correlation function>
Context: colvar
Acceptable values: positive integer
Default value: 1
Description: Number of steps between two values of the time correlation function.

- **corrFuncOffset** 〈Offset of the time correlation function〉
Context: colvar
Acceptable values: positive integer
Default value: 0
Description: The starting time (in number of steps) of the time correlation function (default: $t = 0$).
Note: *the value at $t = 0$ is always used for the normalization.*
- **corrFuncOutputFile** 〈Output file for the time correlation function〉
Context: colvar
Acceptable values: UNIX filename
Default value: <name>.corrfunc.dat
Description: The time correlation function is saved in this file.
- **runAve** 〈Calculate the running average and standard deviation〉
Context: colvar
Acceptable values: boolean
Default value: off
Description: Whether or not the running average and standard deviation should be calculated for this colvar.
- **runAveLength** 〈Length of the running average window〉
Context: colvar
Acceptable values: positive integer
Default value: 1000
Description: Length (in number of points) of the running average window.
- **runAveStride** 〈Stride of the running average window values〉
Context: colvar
Acceptable values: positive integer
Default value: 1
Description: Number of steps between two values within the running average window.
- **runAveOutputFile** 〈Output file for the running average and standard deviation〉
Context: colvar
Acceptable values: UNIX filename
Default value: <name>.runave.dat
Description: The running average and standard deviation are saved in this file.

4 Selecting atoms for colvars: defining atom groups

4.1 Selection keywords

To define collective variables, atoms are usually selected by group. Each group is identified by a name that is unique in the context of the specific colvar component (e.g. for a distance component, the names of the two groups are group1 and group2). The name is followed by a brace-delimited block of selection keywords: these may be used individually or in combination with each other, and each can be repeated any number of times. Selection is incremental: each keyword adds the corresponding atoms to the selection, so that different sets of atoms can be combined. However, atoms included by multiple keywords are only counted once. Below is a example configuration for an atom group with name “atoms”, which uses an unusually varied combination of selection keywords:

```
atoms {

    # add atoms 1 and 3 to this group (note: the first atom in the system is 1)
    atomNumbers {
        1 3
    }

    # add atoms starting from 20 up to and including 50
    atomNumbersRange 20-50

    # add index group (requires a .ndx file to be provided globally)
    indexGroup Water

}
```

The resulting selection includes atoms 1 and 3, those between 20 and 50, and those in the index group called “Water”; the indices of this group are read from the file provided by `indexFile`, in the global section of the configuration file.

The list of selection keywords is:

- `atomNumbers` ⟨List of atom numbers⟩
Context: atom group
Acceptable values: space-separated list of positive integers
Description: This option adds to the group all the atoms whose numbers are in the list. *The number of the first atom in the system is 1: to convert from a VMD selection, use “atomselect get serial”.*
- `indexGroup` ⟨Name of index group to be used (GROMACS format)⟩
Context: atom group
Acceptable values: string
Description: If the name of an index file has been provided by `indexFile`, this option allows to select one index group from that file: the atoms from that index group will be used to define the current group.
- `atomNumbersRange` ⟨Atoms within a number range⟩
Context: atom group
Acceptable values: <Starting number>-<Ending number>
Description: This option includes in the group all atoms whose numbers are within the range specified. *The number of the first atom in the system is 1.*
- `dummyAtom` ⟨Dummy atom position (Å)⟩
Context: atom group
Acceptable values: (x, y, z) triplet
Description: Instead of selecting any atom, this option makes the group a virtual particle at a fixed position in space. This is useful e.g. to replace a group’s center of geometry with a user-defined position.

4.2 Moving frame of reference.

The following options define an automatic calculation of an optimal translation (centerReference) or optimal rotation (rotateReference), that superimposes the positions of this group to a provided set of reference coordinates. This can allow, for example, to effectively remove from certain colvars the effects of molecular tumbling and of diffusion. Given the set of atomic positions \mathbf{x}_i , the colvar ξ can be defined on a set of roto-translated positions $\mathbf{x}'_i = R(\mathbf{x}_i - \mathbf{x}^C) + \mathbf{x}^{\text{ref}}$. \mathbf{x}^C is the geometric center of the \mathbf{x}_i , R is the optimal rotation matrix to the reference positions and \mathbf{x}^{ref} is the geometric center of the reference positions.

Components that are defined based on pairwise distances are naturally invariant under global roto-translations. Other components are instead affected by global rotations or translations: however, they can be made invariant if they are expressed in the frame of reference of a chosen group of atoms, using the centerReference and rotateReference options. Finally, a few components are defined by convention using a roto-translated frame (e.g. the minimal RMSD): for these components, centerReference and rotateReference are enabled by default. In typical applications, the default settings result in the expected behavior.

- centerReference \langle Implicitly remove translations for this group \rangle
Context: atom group
Acceptable values: boolean
Default value: off
Description: If this option is on, the center of geometry of the group will be aligned with that of the reference positions provided by refPositions. Colvar components will only have access to the aligned positions. **Note:** unless otherwise specified, rmsd and eigenvector set this option to on *by default*.
- rotateReference \langle Implicitly remove rotations for this group \rangle
Context: atom group
Acceptable values: boolean
Default value: off
Description: If this option is on, the coordinates of this group will be optimally superimposed to the reference positions provided by refPositions. The rotation will be performed around the center of geometry if centerReference is on, around the origin otherwise. The algorithm used is the same employed by the orientation colvar component [2]. Forces applied to the atoms of this group will also be implicitly rotated back to the original frame. **Note:** unless otherwise specified, rmsd and eigenvector set this option to on *by default*.
- refPositions \langle Reference positions for fitting (\AA) \rangle
Context: atom group
Acceptable values: space-separated list of (x, y, z) triplets
Description: This option provides a list of reference coordinates for centerReference or rotateReference. If only centerReference is on, the list may contain a single (x, y, z) triplet; if also rotateReference is on, the list should be as long as the atom group.
- refPositionsGroup \langle Use an alternate set of atoms to define the roto-translation \rangle
Context: atom group
Acceptable values: Block refPositionsGroup { ... }
Default value: This group itself
Description: If either centerReference or rotateReference is defined, this keyword defines an alternate atom group to calculate the optimal roto-translation. Use this option to define a continuous

rotation if the structure of the group involved changes significantly (a typical symptom would be the message “Warning: discontinuous rotation!”).

The following two options have default values appropriate for the vast majority of applications, and are only provided to support rare, special cases.

- **enableFitGradients** 〈Include the roto-translational contribution to colvar gradients〉
Context: atom group
Acceptable values: boolean
Default value: on
Description: When either `centerReference` or `rotateReference` is on, the gradients of some colvars include terms proportional to $\partial R / \partial \mathbf{x}_i$ (rotational gradients) and $\partial \mathbf{x}^C / \partial \mathbf{x}_i$ (translational gradients). By default, these terms are calculated and included in the total gradients; if this option is set to off, they are neglected.
- **enableForces** 〈Apply forces from this colvar to this group〉
Context: atom group
Acceptable values: boolean
Default value: on
Description: If this option is off, no forces are applied from this colvar to this group. Other forces are not affected (i.e. those from the MD engine, from other colvars, and other external forces). For dummy atoms, this option is on by default.

4.3 Treatment of periodic boundary conditions.

In simulations with periodic boundary conditions, many of the implemented colvar components rely on the fact that each position within a group of atoms is at the nearest periodic image from the center of geometry of the group itself. However, due to the internal wrapping of individual atomic positions done by LAMMPS, this assumption is inaccurate if groups lie at one of the unit cell’s boundaries. For this reason, within the colvars module coordinates are unwrapped by default to avoid discontinuities due to coordinate wrapping (see `unwrap` keyword in 2.1). The user should determine whether maintaining the default value of `unwrap`, depending on the specifics of each system. In general, internal coordinate wrapping by LAMMPS does not affect the calculation of colvars if each atom group satisfies one or more of the following:

- it is composed by only one atom;
- it is used by a colvar component which does not make use of its center of geometry, but only of pairwise distances (`distanceInv`, `coordNum`, `hBond`, `alpha`, `dihedralPC`);
- it is used by a colvar component that ignores the ill-defined Cartesian components of its center of mass (such as the x and y components of a membrane’s center of mass modeled with `distanceZ`).

4.4 Computational cost of colvars based on group size.

In parallel MD simulations, the calculation of most interaction terms are spread over many computational nodes, but the calculation of colvars is not parallelized. Therefore, additional calculations are executed by the node calculating the colvars, and most importantly, additional communication is added between the first node and the other nodes. To mitigate that, atom groups should be kept relatively small (up to a few thousands, depending on the computational cost to simulate the system by itself).

5 Collective variable components (basis functions)

Each colvar is defined by one or more *components* (typically only one). Each component consists of a keyword identifying a functional form, and a definition block following that keyword, specifying the atoms involved and any additional parameters (cutoffs, “reference” values, ...).

The types of the components used in a colvar determine the properties of that colvar, and which biasing or analysis methods can be applied. In most cases, the colvar returns a real number, which is computed by one or more instances of the following components:

- distance: distance between two groups;
- distanceZ: projection of a distance vector on an axis;
- distanceXY: projection of a distance vector on a plane;
- distanceVec: distance vector between two groups;
- distanceDir: unit vector parallel to distanceVec;
- distanceInv: mean distance between two groups of atoms (e.g. NOE-based distance);
- angle: angle between three groups;
- coordNum: coordination number between two groups;
- selfCoordNum: coordination number of atoms within a group;
- hBond: hydrogen bond between two atoms;
- rmsd: root mean square deviation (RMSD) from a set of reference coordinates;
- eigenvector: projection of the atomic coordinates on a vector;
- orientationAngle: angle of the best-fit rotation from a set of reference coordinates;
- spinAngle: projection orthogonal to an axis of the best-fit rotation from a set of reference coordinates;
- tilt: projection on an axis of the best-fit rotation from a set of reference coordinates;
- gyration: radius of gyration of a group of atoms;
- inertia: moment of inertia of a group of atoms;
- inertiaZ: moment of inertia of a group of atoms around a chosen axis;

In the following, all the available component types are listed, along with their physical units and the limiting values, if any. Such limiting values can be used to define `lowerBoundary` and `upperBoundary` in the parent colvar.

5.1 List of available colvar components

5.1.1 distance: center-of-mass distance between two groups.

The distance `{...}` block defines a distance component, between two atom groups, group1 and group2.

- `group1` `<First group of atoms>`
Context: distance
Acceptable values: Block `group1 {...}`
Description: First group of atoms.
- `group2` `<Second group of atoms>`
Context: distance
Acceptable values: Block `group2 {...}`
Description: Second group of atoms.
- `forceNoPBC` `<Calculate absolute rather than minimum-image distance?>`
Context: distance
Acceptable values: boolean
Default value: no
Description: By default, in calculations with periodic boundary conditions, the distance component returns the distance according to the minimum-image convention. If this parameter is set to yes, PBC will be ignored and the distance between the coordinates as maintained internally will be used. This is only useful in a limited number of special cases, e.g. to describe the distance between remote points of a single macromolecule, which cannot be split across periodic cell boundaries, and for which the minimum-image distance might give the wrong result because of a relatively small periodic cell.
- `oneSiteSystemForce` `<Measure system force on group 1 only?>`
Context: distance
Acceptable values: boolean
Default value: no
Description: If this is set to yes, the system force is measured along a vector field (see equation (15) in section 6.1) that only involves atoms of group1. This option is only useful for ABF, or custom biases that compute system forces. See section 6.1 for details.

The value returned is a positive number (in Å), ranging from 0 to the largest possible interatomic distance within the chosen boundary conditions (with PBCs, the minimum image convention is used unless the `forceNoPBC` option is set).

5.1.2 distanceZ: projection of a distance vector on an axis.

The `distanceZ {...}` block defines a distance projection component, which can be seen as measuring the distance between two groups projected onto an axis, or the position of a group along such an axis. The axis can be defined using either one reference group and a constant vector, or dynamically based on two reference groups.

- `main` `<Main group of atoms>`
Context: `distanceZ`, `distanceXY`
Acceptable values: Block `main {...}`
Description: Group of atoms whose position r is measured.

- **ref** <Reference group of atoms>
Context: distanceZ, distanceXY
Acceptable values: Block ref {...}
Description: Reference group of atoms. The position of its center of mass is noted r_1 below.
- **ref2** <Secondary reference group>
Context: distanceZ, distanceXY
Acceptable values: Block ref2 {...}
Default value: none
Description: Optional group of reference atoms, whose position r_2 can be used to define a dynamic projection axis: $e = (\|r_2 - r_1\|)^{-1} \times (r_2 - r_1)$. In this case, the origin is $r_m = 1/2(r_1 + r_2)$, and the value of the component is $e \cdot (r - r_m)$.
- **axis** <Projection axis (Å)>
Context: distanceZ, distanceXY
Acceptable values: (x, y, z) triplet
Default value: (0.0, 0.0, 1.0)
Description: The three components of this vector define (when normalized) a projection axis e for the distance vector $r - r_1$ joining the centers of groups ref and main. The value of the component is then $e \cdot (r - r_1)$. The vector should be written as three components separated by commas and enclosed in parentheses.
- **forceNoPBC** <Calculate absolute rather than minimum-image distance?>
Context: distanceZ, distanceXY
Acceptable values: boolean
Default value: no
Description: This parameter has the same meaning as that described above for the distance component.
- **oneSiteSystemForce** <Measure system force on group main only?>
Context: distanceZ, distanceXY
Acceptable values: boolean
Default value: no
Description: If this is set to yes, the system force is measured along a vector field (see equation (15) in section 6.1) that only involves atoms of main. This option is only useful for ABF, or custom biases that compute system forces. See section 6.1 for details.

This component returns a number (in Å) whose range is determined by the chosen boundary conditions. For instance, if the z axis is used in a simulation with periodic boundaries, the returned value ranges between $-b_z/2$ and $b_z/2$, where b_z is the box length along z (this behavior is disabled if forceNoPBC is set).

5.1.3 distanceXY: modulus of the projection of a distance vector on a plane.

The distanceXY {...} block defines a distance projected on a plane, and accepts the same keywords as the component distanceZ, i.e. main, ref, either ref2 or axis, and oneSiteSystemForce. It returns the norm of the projection of the distance vector between main and ref onto the plane orthogonal to the axis. The axis is defined using the axis parameter or as the vector joining ref and ref2 (see distanceZ above).

5.1.4 distanceVec: distance vector between two groups.

The distanceVec {...} block defines a distance vector component, which accepts the same keywords as the component distance: group1, group2, and forceNoPBC. Its value is the 3-vector joining the centers of mass of group1 and group2.

5.1.5 distanceDir: distance unit vector between two groups.

The distanceDir {...} block defines a distance unit vector component, which accepts the same keywords as the component distance: group1, group2, and forceNoPBC. It returns a 3-dimensional unit vector $\mathbf{d} = (d_x, d_y, d_z)$, with $|\mathbf{d}| = 1$.

5.1.6 distanceInv: mean distance between two groups of atoms.

The distanceInv {...} block defines a generalized mean distance between two groups of atoms 1 and 2, weighted with exponent $1/n$:

$$d_{1,2}^{[n]} = \left(\frac{1}{N_1 N_2} \sum_{i,j} \left(\frac{1}{\|\mathbf{d}^{ij}\|} \right)^n \right)^{-1/n} \quad (2)$$

where $\|\mathbf{d}^{ij}\|$ is the distance between atoms i and j in groups 1 and 2 respectively, and n is an even integer. This component accepts the same keywords as the component distance: group1, group2, and forceNoPBC. In addition, the following option may be provided:

- exponent <Exponent n in equation 2>

Context: distanceInv

Acceptable values: positive even integer

Default value: 6

Description: Defines the exponent to which the individual distances are elevated before averaging. The default value of 6 is useful for example to applying restraints based on NOE-measured distances.

This component returns a number in Å, ranging from 0 to the largest possible distance within the chosen boundary conditions.

5.1.7 angle: angle between three groups.

The angle {...} block defines an angle, and contains the three blocks group1, group2 and group3, defining the three groups. It returns an angle (in degrees) within the interval $[0 : 180]$.

5.1.8 dihedral: torsional angle between four groups.

The dihedral {...} block defines a torsional angle, and contains the blocks group1, group2, group3 and group4, defining the four groups. It returns an angle (in degrees) within the interval $[-180 : 180]$. The colvar module calculates all the distances between two angles taking into account periodicity. For instance, reference values for restraints or range boundaries can be defined by using any real number of choice.

- oneSiteSystemForce <Measure system force on group 1 only?>

Context: angle, dihedral

Acceptable values: boolean

Default value: no

Description: If this is set to yes, the system force is measured along a vector field (see equation (15) in section 6.1) that only involves atoms of group1. See section 6.1 for an example.

5.1.9 coordNum: coordination number between two groups.

The coordNum { ... } block defines a coordination number (or number of contacts), which calculates the function $(1 - (d/d_0)^n)/(1 - (d/d_0)^m)$, where d_0 is the “cutoff” distance, and n and m are exponents that can control its long range behavior and stiffness [1]. This function is summed over all pairs of atoms in group1 and group2:

$$C(\text{group1}, \text{group2}) = \sum_{i \in \text{group1}} \sum_{j \in \text{group2}} \frac{1 - (|\mathbf{x}_i - \mathbf{x}_j|/d_0)^n}{1 - (|\mathbf{x}_i - \mathbf{x}_j|/d_0)^m} \quad (3)$$

This colvar component accepts the same keywords as the component distance, group1 and group2. In addition to them, it recognizes the following keywords:

- **cutoff** <“Interaction” distance (Å)>
Context: coordNum
Acceptable values: positive decimal
Default value: 4.0
Description: This number defines the switching distance to define an interatomic contact: for $d \ll d_0$, the switching function $(1 - (d/d_0)^n)/(1 - (d/d_0)^m)$ is close to 1, at $d = d_0$ it has a value of n/m ($1/2$ with the default n and m), and at $d \gg d_0$ it goes to zero approximately like d^{m-n} . Hence, for a proper behavior, m must be larger than n .
- **cutoff3** <Reference distance vector (Å)>
Context: coordNum
Acceptable values: “(x, y, z)” triplet of positive decimals
Default value: (4.0, 4.0, 4.0)
Description: The three components of this vector define three different cutoffs d_0 for each direction. This option is mutually exclusive with cutoff.
- **expNumer** <Numerator exponent>
Context: coordNum
Acceptable values: positive even integer
Default value: 6
Description: This number defines the n exponent for the switching function.
- **expDenom** <Denominator exponent>
Context: coordNum
Acceptable values: positive even integer
Default value: 12
Description: This number defines the m exponent for the switching function.
- **group2CenterOnly** <Use only group2’s center of mass>
Context: coordNum
Acceptable values: boolean
Default value: off
Description: If this option is on, only contacts between the atoms in group1 and the center of mass of group2 are calculated. By default, the sum extends over all pairs of atoms in group1 and group2.

This component returns a dimensionless number, which ranges from approximately 0 (all interatomic distances are much larger than the cutoff) to $N_{\text{group1}} \times N_{\text{group2}}$ (all distances are less than the cutoff), or N_{group1} if group2CenterOnly is used. For performance reasons, at least one of group1 and group2 should be of limited size or group2CenterOnly should be used: the cost of the loop over all pairs grows as $N_{\text{group1}} \times N_{\text{group2}}$.

5.1.10 selfCoordNum: coordination number between atoms within a group.

The selfCoordNum $\{\dots\}$ block defines a coordination number similarly to the component coordNum, but the function is summed over atom pairs within group1:

$$C(\text{group1}) = \sum_{i \in \text{group1}} \sum_{j > i} \frac{1 - (|\mathbf{x}_i - \mathbf{x}_j|/d_0)^n}{1 - (|\mathbf{x}_i - \mathbf{x}_j|/d_0)^m} \quad (4)$$

The keywords accepted by selfCoordNum are a subset of those accepted by coordNum, namely group1 (here defining *all* of the atoms to be considered), cutoff, expNumer, and expDenom.

This component returns a dimensionless number, which ranges from approximately 0 (all interatomic distances much larger than the cutoff) to $N_{\text{group1}} \times (N_{\text{group1}} - 1)/2$ (all distances within the cutoff). For performance reasons, group1 should be of limited size, because the cost of the loop over all pairs grows as N_{group1}^2 .

hBond: hydrogen bond between two atoms. The hBond $\{\dots\}$ block defines a hydrogen bond, implemented as a coordination number (eq. 3) between the donor and the acceptor atoms. Therefore, it accepts the same options cutoff (with a different default value of 3.3 Å), expNumer (with a default value of 6) and expDenom (with a default value of 8). Unlike coordNum, it requires two atom numbers, acceptor and donor, to be defined. It returns an adimensional number, with values between 0 (acceptor and donor far outside the cutoff distance) and 1 (acceptor and donor much closer than the cutoff).

5.1.11 rmsd: root mean square displacement (RMSD) from reference positions.

The block rmsd $\{\dots\}$ defines the root mean square replacement (RMSD) of a group of atoms with respect to a reference structure. For each set of coordinates $\{\mathbf{x}_1(t), \mathbf{x}_2(t), \dots, \mathbf{x}_N(t)\}$, the colvar component rmsd calculates the optimal rotation $U^{\{\mathbf{x}_i(t)\} \rightarrow \{\mathbf{x}_i^{(\text{ref})}\}}$ that best superimposes the coordinates $\{\mathbf{x}_i(t)\}$ onto a set of reference coordinates $\{\mathbf{x}_i^{(\text{ref})}\}$. Both the current and the reference coordinates are centered on their centers of geometry, $\mathbf{x}_{\text{cog}}(t)$ and $\mathbf{x}_{\text{cog}}^{(\text{ref})}$. The root mean square displacement is then defined as:

$$\text{RMSD}(\{\mathbf{x}_i(t)\}, \{\mathbf{x}_i^{(\text{ref})}\}) = \sqrt{\frac{1}{N} \sum_{i=1}^N \left| U(\mathbf{x}_i(t) - \mathbf{x}_{\text{cog}}(t)) - (\mathbf{x}_i^{(\text{ref})} - \mathbf{x}_{\text{cog}}^{(\text{ref})}) \right|^2} \quad (5)$$

The optimal rotation $U^{\{\mathbf{x}_i(t)\} \rightarrow \{\mathbf{x}_i^{(\text{ref})}\}}$ is calculated within the formalism developed in reference [2], which guarantees a continuous dependence of $U^{\{\mathbf{x}_i(t)\} \rightarrow \{\mathbf{x}_i^{(\text{ref})}\}}$ with respect to $\{\mathbf{x}_i(t)\}$. The options for rmsd are:

- atoms $\langle \text{Atom group} \rangle$
Context: rmsd
Acceptable values: atoms $\{\dots\}$ block
Description: Defines the group of atoms of which the RMSD should be calculated. Optimal fit options (such as refPositions and rotateReference) should typically NOT be set within this block. Exceptions to this rule are the special cases discussed in the *Advanced usage* paragraph below.
- refPositions $\langle \text{Reference coordinates} \rangle$
Context: rmsd
Acceptable values: space-separated list of (x, y, z) triplets
Description: This option sets the reference coordinates. If only centerReference is on, the list can be a single (x, y, z) triplet; if also rotateReference is on, the list should be as long as the atom group. This option is independent from that with the same keyword within the atoms $\{\dots\}$ block.

The latter (and related fitting options for the atom group) are normally not needed, and should be omitted altogether except for advanced usage cases.

This component returns a positive real number (in Å).

5.1.12 Advanced usage of the rmsd component.

In the standard usage as described above, the rmsd component calculates a minimum RMSD, that is, current coordinates are optimally fitted onto the same reference coordinates that are used to compute the RMSD value. The fit itself is handled by the atom group object, whose parameters are automatically set by the rmsd component. For very specific applications, however, it may be useful to control the fitting process separately from the definition of the reference coordinates, to evaluate various types of non-minimal RMSD values. This can be achieved by setting the related options (refPositions, etc.) explicitly in the atom group block. This allows for the following non-standard cases:

1. applying the optimal translation, but no rotation (rotateReference off), to bias or restrain the shape and orientation, but not the position of the atom group;
2. applying the optimal rotation, but no translation (translateReference off), to bias or restrain the shape and position, but not the orientation of the atom group;
3. disabling the application of optimal roto-translations, which lets the RMSD component describe the deviation of atoms from fixed positions in the laboratory frame: this allows for custom positional restraints within the colvars module;
4. fitting the atomic positions to different reference coordinates than those used in the RMSD calculation itself;
5. applying the optimal rotation and/or translation from a separate atom group, defined through refPositionsGroup: the RMSD then reflects the deviation from reference coordinates in a separate, moving reference frame.

5.1.13 eigenvector: projection of the atomic coordinates on a vector.

The block eigenvector $\{\dots\}$ defines the projection of the coordinates of a group of atoms (or more precisely, their deviations from the reference coordinates) onto a vector in \mathbb{R}^{3n} , where n is the number of atoms in the group. The computed quantity is the total projection:

$$p(\{\mathbf{x}_i(t)\}, \{\mathbf{x}_i^{(\text{ref})}\}) = \sum_{i=1}^n \mathbf{v}_i \cdot \left(U(\mathbf{x}_i(t) - \mathbf{x}_{\text{cog}}(t)) - (\mathbf{x}_i^{(\text{ref})} - \mathbf{x}_{\text{cog}}^{(\text{ref})}) \right), \quad (6)$$

where, as in the rmsd component, U is the optimal rotation matrix, $\mathbf{x}_{\text{cog}}(t)$ and $\mathbf{x}_{\text{cog}}^{(\text{ref})}$ are the centers of geometry of the current and reference positions respectively, and \mathbf{v}_i are the components of the vector for each atom. Example choices for (\mathbf{v}_i) are an eigenvector of the covariance matrix (essential mode), or a normal mode of the system. It is assumed that $\sum_i \mathbf{v}_i = 0$: otherwise, the colvars module centers the \mathbf{v}_i automatically when reading them from the configuration.

As in the rmsd component, available options are atoms and refPositions. In addition, the following are recognized:

- **vector** \langle Vector components \rangle
Context: eigenvector
Acceptable values: space-separated list of (x, y, z) triplets
Description: This option sets the values of the vector components.
- **differenceVector** \langle The $3n$ -dimensional vector is the difference between vector and refPositions \rangle
Context: eigenvector
Acceptable values: boolean
Default value: off
Description: If this option is on, the numbers provided by vector are interpreted as another set of positions, \mathbf{x}'_i : the vector \mathbf{v}_i is then defined as $\mathbf{v}_i = (\mathbf{x}'_i - \mathbf{x}_i^{(\text{ref})})$. This allows to conveniently define a colvar ξ as a projection on the linear transformation between two sets of positions, “A” and “B”. For convenience, the vector is also normalized so that $\xi = 0$ when the atoms are at the set of positions “A” and $\xi = 1$ at the set of positions “B”.

This component returns a number (in Å), whose value ranges between the smallest and largest absolute positions in the unit cell during the simulations (see also distanceZ). Due to the normalization in eq. 6, this range does not depend on the number of atoms involved.

5.1.14 gyration: radius of gyration of a group of atoms.

The block gyration $\{\dots\}$ defines the parameters for calculating the radius of gyration of a group of atomic positions $\{\mathbf{x}_1(t), \mathbf{x}_2(t), \dots, \mathbf{x}_N(t)\}$ with respect to their center of geometry, $\mathbf{x}_{\text{cog}}(t)$:

$$R_{\text{gyr}} = \sqrt{\frac{1}{N} \sum_{i=1}^N |\mathbf{x}_i(t) - \mathbf{x}_{\text{cog}}(t)|^2} \quad (7)$$

This component must contain one atoms $\{\dots\}$ block to define the atom group, and returns a positive number, expressed in Å.

5.1.15 inertia: total moment of inertia of a group of atoms.

The block inertia $\{\dots\}$ defines the parameters for calculating the total moment of inertia of a group of atomic positions $\{\mathbf{x}_1(t), \mathbf{x}_2(t), \dots, \mathbf{x}_N(t)\}$ with respect to their center of geometry, $\mathbf{x}_{\text{cog}}(t)$:

$$I = \sum_{i=1}^N |\mathbf{x}_i(t) - \mathbf{x}_{\text{cog}}(t)|^2 \quad (8)$$

Note that all atomic masses are set to 1 for simplicity. This component must contain one atoms $\{\dots\}$ block to define the atom group, and returns a positive number, expressed in Å².

5.1.16 inertiaZ: total moment of inertia of a group of atoms around a chosen axis.

The block inertiaZ $\{\dots\}$ defines the parameters for calculating the component along the axis \mathbf{e} of the moment of inertia of a group of atomic positions $\{\mathbf{x}_1(t), \mathbf{x}_2(t), \dots, \mathbf{x}_N(t)\}$ with respect to their center of geometry, $\mathbf{x}_{\text{cog}}(t)$:

$$I_{\mathbf{e}} = \sum_{i=1}^N ((\mathbf{x}_i(t) - \mathbf{x}_{\text{cog}}(t)) \cdot \mathbf{e})^2 \quad (9)$$

Note that all atomic masses are set to 1 for simplicity. This component must contain one `atoms { ... }` block to define the atom group, and returns a positive number, expressed in \AA^2 . The following option may also be provided:

- `axis` \langle Projection axis (\AA) \rangle
Context: `inertiaZ`
Acceptable values: (x, y, z) triplet
Default value: (0.0, 0.0, 1.0)
Description: The three components of this vector define (when normalized) the projection axis **e**.

5.1.17 orientation: orientation from reference coordinates.

The block `orientation { ... }` returns the same optimal rotation used in the `rmsd` component to superimpose the coordinates $\{\mathbf{x}_i(t)\}$ onto a set of reference coordinates $\{\mathbf{x}_i^{(\text{ref})}\}$. Such component returns a four dimensional vector $\mathbf{q} = (q_0, q_1, q_2, q_3)$, with $\sum_i q_i^2 = 1$; this *quaternion* expresses the optimal rotation $\{\mathbf{x}_i(t)\} \rightarrow \{\mathbf{x}_i^{(\text{ref})}\}$ according to the formalism in reference [2]. The quaternion (q_0, q_1, q_2, q_3) can also be written as $(\cos(\theta/2), \sin(\theta/2)\mathbf{u})$, where θ is the angle and **u** the normalized axis of rotation; for example, a rotation of 90° around the *z* axis should be expressed as “(0.707, 0.0, 0.0, 0.707)”. The script `quaternion2rmatrix.tcl` provides Tcl functions for converting to and from a 4×4 rotation matrix in a format suitable for usage in VMD.

As for the component `rmsd`, the options `atoms` and `refPositions` are supported. **Note:** `refPositions` define the set of positions *from which* the optimal rotation is calculated, but this rotation is not applied to the coordinates of the atoms involved: it is used instead to define the variable itself.

- `closestToQuaternion` \langle Reference rotation \rangle
Context: `orientation`
Acceptable values: “(q0, q1, q2, q3)” quadruplet
Default value: (1.0, 0.0, 0.0, 0.0) (“null” rotation)
Description: Between the two equivalent quaternions (q_0, q_1, q_2, q_3) and $(-q_0, -q_1, -q_2, -q_3)$, the closer to (1.0, 0.0, 0.0, 0.0) is chosen. This simplifies the visualization of the colvar trajectory when samples values are a smaller subset of all possible rotations. **Note:** *this only affects the output, never the dynamics.*

Hint: stopping the rotation of a protein. To stop the rotation of an elongated macromolecule in solution (and use an anisotropic box to save water molecules), it is possible to define a colvar with an orientation component, and restrain it through the harmonic bias around the identity rotation, (1.0, 0.0, 0.0, 0.0). Only the overall orientation of the macromolecule is affected, and *not* its internal degrees of freedom. The user should also take care that the macromolecule is composed by a single chain, or disable `wrapAll` otherwise.

5.1.18 orientationAngle: angle of rotation from reference coordinates.

The block `orientationAngle { ... }` accepts the same base options as the component `orientation`: `atoms`, `refPositions`. The returned value is the angle of rotation ω between the current and the reference positions. This angle is expressed in degrees within the range $[0^\circ:180^\circ]$.

5.1.19 spinAngle: angle of rotation around a given axis.

The complete rotation described by `orientation` can optionally be decomposed into two sub-rotations: one is a “*spin*” rotation around **e**, and the other a “*tilt*” rotation around an axis orthogonal to **e**. The component

spinAngle measures the angle of the “spin” sub-rotation around **e**. This can be defined using the same options as the component orientation: atoms, refPositions. In addition, spinAngle accepts the axis option:

- axis <Special rotation axis (Å)>

Context: tilt, spinAngle

Acceptable values: (x, y, z) triplet

Default value: (0.0, 0.0, 1.0)

Description: The three components of this vector define (when normalized) the special rotation axis used to calculate the tilt and spinAngle components.

The component spinAngle returns an angle (in degrees) within the periodic interval $[-180 : 180]$.

Note: the value of spinAngle is a continuous function almost everywhere, with the exception of configurations with the corresponding “tilt” angle equal to 180° (i.e. the tilt component is equal to -1): in those cases, spinAngle is undefined. If such configurations are expected, consider defining a tilt colvar using the same axis **e**, and restraining it with a lower wall away from -1 .

5.1.20 tilt: cosine of the rotation orthogonal to a given axis.

The component tilt measures the cosine of the angle of the “tilt” sub-rotation, which combined with the “spin” sub-rotation provides the complete rotation of a group of atoms. The cosine of the tilt angle rather than the tilt angle itself is implemented, because the latter is unevenly distributed even for an isotropic system: consider as an analogy the angle θ in the spherical coordinate system. The component tilt relies on the same options as spinAngle, including the definition of the axis **e**. The values of tilt are real numbers in the interval $[-1 : 1]$: the value 1 represents an orientation fully parallel to **e** (tilt angle = 0°), and the value -1 represents an anti-parallel orientation.

5.2 Advanced usage and special considerations

5.2.1 Periodic components.

The following components returns real numbers that lie in a periodic interval:

- dihedral: torsional angle between four groups;
- spinAngle: angle of rotation around a predefined axis in the best-fit from a set of reference coordinates.

In certain conditions, distanceZ can also be periodic, namely when periodic boundary conditions (PBCs) are defined in the simulation and distanceZ’s axis is parallel to a unit cell vector.

The following keywords can be used within periodic components (and are illegal elsewhere):

- period <Period of the component>

Context: distanceZ

Acceptable values: positive decimal

Default value: 0.0

Description: Setting this number enables the treatment of distanceZ as a periodic component: by default, distanceZ is not considered periodic. The keyword is supported, but irrelevant within dihedral or spinAngle, because their period is always 360 degrees.

- `wrapAround` <Center of the wrapping interval for periodic variables>

Context: `distanceZ`, `dihedral` or `spinAngle`

Acceptable values: decimal

Default value: 0.0

Description: By default, values of the periodic components are centered around zero, ranging from $-P/2$ to $P/2$, where P is the period. Setting this number centers the interval around this value. This can be useful for convenience of output, or to set `lowerWall` and `upperWall` in an order that would not otherwise be allowed.

Internally, all differences between two values of a periodic colvar follow the minimum image convention: they are calculated based on the two periodic images that are closest to each other.

Note: linear or polynomial combinations of periodic components may become meaningless when components cross the periodic boundary. Use such combinations carefully: estimate the range of possible values of each component in a given simulation, and make use of `wrapAround` to limit this problem whenever possible.

5.2.2 Non-scalar components.

When one of the following are used, the colvar returns a value that is not a scalar number:

- `distanceVec`: 3-dimensional vector of the distance between two groups;
- `distanceDir`: 3-dimensional unit vector of the distance between two groups;
- `orientation`: 4-dimensional unit quaternion representing the best-fit rotation from a set of reference coordinates.

The distance between two 3-dimensional unit vectors is computed as the angle between them. The distance between two quaternions is computed as the angle between the two 4-dimensional unit vectors: because the orientation represented by q is the same as the one represented by $-q$, distances between two quaternions are computed considering the closest of the two symmetric images.

Non-scalar components carry the following restrictions:

- Calculation of system forces (`outputSystemForce` option) is currently not implemented.
- Each colvar can only contain one non-scalar component.
- Binning on a grid (`abf`, `histogram` and `metadynamics` with `useGrids` enabled) is currently not implemented for colvars based on such components.

Note: while these restrictions apply to individual colvars based on non-scalar components, no limit is set to the number of scalar colvars. To compute multi-dimensional histograms and PMFs, use sets of scalar colvars of arbitrary size.

5.2.3 Calculating system forces.

In addition to the restrictions due to the type of value computed (scalar or non-scalar), a final restriction can arise when calculating system force (`outputSystemForce` option or application of a `abf` bias). System forces are available currently only for the following components: `distance`, `distanceZ`, `distanceXY`, `angle`, `dihedral`, `rmsd`, `eigenvector` and `gyration`.

5.3 Linear and polynomial combinations of components

To extend the set of possible definitions of colvars $\xi(\mathbf{r})$, multiple components $q_i(\mathbf{r})$ can be summed with the formula:

$$\xi(\mathbf{r}) = \sum_i c_i [q_i(\mathbf{r})]^{n_i} \quad (10)$$

where each component appears with a unique coefficient c_i (1.0 by default) the positive integer exponent n_i (1 by default).

Any set of components can be combined within a colvar, provided that they return the same type of values (scalar, unit vector, vector, or quaternion). By default, the colvar is the sum of its components. Linear or polynomial combinations (following equation (10)) can be obtained by setting the following parameters, which are common to all components:

- **componentCoeff** $\langle \text{decimal} \rangle$
Context: (any component) Coefficient of this component in the colvar
Acceptable values: 1.0
Default value: Defines the coefficient by which this component is multiplied (after being raised to componentExp) before being added to the sum.
Description:
- **componentExp** $\langle \text{integer} \rangle$
Context: (any component) Exponent of this component in the colvar
Acceptable values: 1
Default value: Defines the power at which the value of this component is raised before being added to the sum. When this exponent is different than 1 (non-linear sum), system forces and the Jacobian force are not available, making the colvar unsuitable for ABF calculations.
Description:

Example: To define the *average* of a colvar across different parts of the system, simply define within the same colvar block a series of components of the same type (applied to different atom groups), and assign to each component a componentCoeff of $1/N$.

6 Biasing and analysis methods

All of the biasing and analysis methods implemented (abf, harmonic, histogram and metadynamics) recognize the following options:

- **name** $\langle \text{Identifier for the bias} \rangle$
Context: colvar bias
Acceptable values: string
Default value: $\langle \text{type of bias} \rangle \langle \text{bias index} \rangle$
Description: This string is used to identify the bias or analysis method in output messages and to name some output files.
- **colvars** $\langle \text{Collective variables involved} \rangle$
Context: colvar bias
Acceptable values: space-separated list of colvar names
Description: This option selects by name all the colvars to which this bias or analysis will be applied.

- outputEnergy \langle Write the current bias energy to the trajectory file \rangle

Context: colvar bias

Acceptable values: boolean

Default value: off

Description: If this option is chosen and colvarsTrajFrequency is not zero, the current value of the biasing energy will be written to the trajectory file during the simulation.

6.1 Adaptive Biasing Force

For a full description of the Adaptive Biasing Force method, see reference [3]. For details about this implementation, see references [4] and [?]. **When publishing research that makes use of this functionality, please cite references [3] and [?].**

An alternate usage of this feature is the application of custom tabulated biasing potentials to one or more colvars. See inputPrefix and updateBias below.

ABF is based on the thermodynamic integration (TI) scheme for computing free energy profiles. The free energy as a function of a set of collective variables $\xi = (\xi_i)_{i \in [1, n]}$ is defined from the canonical distribution of ξ , $\mathcal{P}(\xi)$:

$$A(\xi) = -\frac{1}{\beta} \ln \mathcal{P}(\xi) + A_0 \quad (11)$$

In the TI formalism, the free energy is obtained from its gradient, which is generally calculated in the form of the average of a force F_ξ exerted on ξ , taken over an iso- ξ surface:

$$\nabla_\xi A(\xi) = \langle -F_\xi \rangle_\xi \quad (12)$$

Several formulae that take the form of (12) have been proposed. This implementation relies partly on the classic formulation [5], and partly on a more versatile scheme originating in a work by Ruiz-Montero et al. [6], generalized by den Otter [7] and extended to multiple variables by Ciccotti et al. [8]. Consider a system subject to constraints of the form $\sigma_k(x) = 0$. Let $(v_i)_{i \in [1, n]}$ be arbitrarily chosen vector fields ($\mathbb{R}^{3N} \rightarrow \mathbb{R}^{3N}$) verifying, for all i, j , and k :

$$v_i \cdot \nabla_x \xi_j = \delta_{ij} \quad (13)$$

$$v_i \cdot \nabla_x \sigma_k = 0 \quad (14)$$

then the following holds [8]:

$$\frac{\partial A}{\partial \xi_i} = \langle v_i \cdot \nabla_x V - k_B T \nabla_x \cdot v_i \rangle_\xi \quad (15)$$

where V is the potential energy function. v_i can be interpreted as the direction along which the force acting on variable ξ_i is measured, whereas the second term in the average corresponds to the geometric entropy contribution that appears as a Jacobian correction in the classic formalism [5]. Condition (13) states that the direction along which the system force on ξ_i is measured is orthogonal to the gradient of ξ_j , which means that the force measured on ξ_i does not act on ξ_j .

Equation (14) implies that constraint forces are orthogonal to the directions along which the free energy gradient is measured, so that the measurement is effectively performed on unconstrained degrees of freedom.

In the framework of ABF, \mathbf{F}_ξ is accumulated in bins of finite size, $\delta\xi$, thereby providing an estimate of the free energy gradient according to equation (12). The biasing force applied along the collective variables to overcome free energy barriers is calculated as:

$$\mathbf{F}^{\text{ABF}} = \nabla_x \tilde{A}(\xi) \quad (16)$$

where $\nabla_x \tilde{A}$ denotes the current estimate of the free energy gradient at the current point ξ in the collective variable subspace.

As sampling of the phase space proceeds, the estimate $\nabla_x \tilde{A}$ is progressively refined. The biasing force introduced in the equations of motion guarantees that in the bin centered around ξ , the forces acting along the selected collective variables average to zero over time. Eventually, as the underlying free energy surface is canceled by the adaptive bias, evolution of the system along ξ is governed mainly by diffusion. Although this implementation of ABF can in principle be used in arbitrary dimension, a higher-dimension collective variable space is likely to result in sampling difficulties. Most commonly, the number of variables is one or two.

6.1.1 ABF requirements on collective variables

1. *Only linear combinations* of colvar components can be used in ABF calculations.
2. *Availability of system forces* is necessary. The following colvar components can be used in ABF calculations: distance, distance_xy, distance_z, angle, dihedral, gyration, rmsd and eigenvector. Atom groups may not be replaced by dummy atoms, unless they are excluded from the force measurement by specifying `oneSiteSystemForce`, if available.
3. *Mutual orthogonality of colvars*. In a multidimensional ABF calculation, equation (13) must be satisfied for any two colvars ξ_i and ξ_j . Various cases fulfill this orthogonality condition:
 - ξ_i and ξ_j are based on non-overlapping sets of atoms.
 - atoms involved in the force measurement on ξ_i do not participate in the definition of ξ_j . This can be obtained using the option `oneSiteSystemForce` of the distance, angle, and dihedral components (example: Ramachandran angles ϕ , ψ).
 - ξ_i and ξ_j are orthogonal by construction. Useful cases are the sum and difference of two components, or distance_z and distance_xy using the same axis.
4. *Mutual orthogonality of components*: when several components are combined into a colvar, it is assumed that their vectors v_i (equation (15)) are mutually orthogonal. The cases described for colvars in the previous paragraph apply.
5. *Orthogonality of colvars and constraints*: equation 14 can be satisfied in two simple ways, if either no constrained atoms are involved in the force measurement (see point 3 above) or pairs of atoms joined by a constrained bond are part of an *atom group* which only intervenes through its center (center of mass or geometric center) in the force measurement. In the latter case, the contributions of the two atoms to the left-hand side of equation 14 cancel out. For example, all atoms of a rigid TIP3P water molecule can safely be included in an atom group used in a distance component.

Parameters for ABF

The following parameters can be set in the ABF configuration block (in addition to generic bias parameters such as colvars):

- `fullSamples` (Number of samples in a bin prior to application of the ABF)
- Context:** abf

Acceptable values: positive integer

Default value: 200

Description: To avoid nonequilibrium effects in the dynamics of the system, due to large fluctuations of the force exerted along the reaction coordinate, ξ , it is recommended to apply the biasing force only after a reasonable estimate of the latter has been obtained.

- **hideJacobian** 〈Remove geometric entropy term from calculated free energy gradient?〉
Context: abf
Acceptable values: boolean
Default value: no
Description: In a few special cases, most notably distance-based variables, an alternate definition of the potential of mean force is traditionally used, which excludes the Jacobian term describing the effect of geometric entropy on the distribution of the variable. This results, for example, in particle-particle potentials of mean force being flat at large separations. Setting this parameter to yes causes the output data to follow that convention, by removing this contribution from the output gradients while applying internally the corresponding correction to ensure uniform sampling. It is not allowed for colvars with multiple components.
- **outputFreq** 〈Frequency (in timesteps) at which ABF data files are refreshed〉
Context: abf
Acceptable values: positive integer
Default value: Colvar module restart frequency
Description: The files containing the free energy gradient estimate and sampling histogram (and the PMF in one-dimensional calculations) are written on disk at the given time interval.
- **historyFreq** 〈Frequency (in timesteps) at which ABF history files are accumulated〉
Context: abf
Acceptable values: positive integer
Default value: 0
Description: If this number is non-zero, the free energy gradient estimate and sampling histogram (and the PMF in one-dimensional calculations) are appended to files on disk at the given time interval. History file names use the same prefix as output files, with “.hist” appended.
- **inputPrefix** 〈Filename prefix for reading ABF data〉
Context: abf
Acceptable values: list of strings
Description: If this parameter is set, for each item in the list, ABF tries to read a gradient and a sampling files named <inputPrefix>.grad and <inputPrefix>.count. This is done at startup and sets the initial state of the ABF algorithm. The data from all provided files is combined appropriately. Also, the grid definition (min and max values, width) need not be the same that for the current run. This command is useful to piece together data from simulations in different regions of collective variable space, or change the colvar boundary values and widths. Note that it is not recommended to use it to switch to a smaller width, as that will leave some bins empty in the finer data grid. This option is NOT compatible with reading the data from a restart file (input keyword of the `fix ID group-ID colvars` command).
- **applyBias** 〈Apply the ABF bias?〉
Context: abf
Acceptable values: boolean
Default value: yes

Description: If this is set to no, the calculation proceeds normally but the adaptive biasing force is not applied. Data is still collected to compute the free energy gradient. This is mostly intended for testing purposes, and should not be used in routine simulations.

- `updateBias` (Update the ABF bias?)

Context: `abf`

Acceptable values: `boolean`

Default value: `yes`

Description: If this is set to no, the initial biasing force (e.g. read from a restart file or through the input keyword) is not updated during the simulation. As a result, a constant bias is applied. This can be used to apply a custom, tabulated biasing potential to any combination of colvars. To that effect, one should prepare a gradient file containing the gradient of the potential to be applied (negative of the bias force), and a count file containing only values greater than `fullSamples`. These files must match the grid parameters of the colvars.

ABF also depends on parameters from collective variables to define the grid on which free energy gradients are computed. In the direction of each colvar, the grid ranges from `lowerBoundary` to `upperBoundary`, and the bin width (grid spacing) is set by the `width` parameter.

6.1.2 Output files

The ABF bias produces the following files, all in multicolumn ASCII format:

- `output.grad`: current estimate of the free energy gradient (grid), in multicolumn;
- `output.count`: total number of samples collected, on the same grid;
- `output.pmf`: only for one-dimensional calculations, integrated free energy profile or PMF.

If several ABF biases are defined concurrently, their name is inserted to produce unique filenames for output, as in `output.abf1.grad`. This should not be done routinely and could lead to meaningless results: only do it if you know what you are doing!

If the colvar space has been partitioned into sections (*windows*) in which independent ABF simulations have been run, the resulting data can be merged using the the input keyword of the `fix ID group-ID colvars` command.

6.1.3 Post-processing: reconstructing a multidimensional free energy surface

If a one-dimensional calculation is performed, the estimated free energy gradient is automatically integrated and a potential of mean force is written under the file name `outputName.pmf`, in a plain text format that can be read by most data plotting and analysis programs (e.g. `gnuplot`).

In dimension 2 or greater, integrating the discretized gradient becomes non-trivial. The standalone utility `abf.integrate` is provided to perform that task. `abf.integrate` reads the gradient data and uses it to perform a Monte-Carlo (M-C) simulation in discretized collective variable space (specifically, on the same grid used by ABF to discretize the free energy gradient). By default, a history-dependent bias (similar in spirit to metadynamics) is used: at each M-C step, the bias at the current position is incremented by a preset amount (the *hill height*). Upon convergence, this bias counteracts optimally the underlying gradient; it is negated to obtain the estimate of the free energy surface.

`abf.integrate` is invoked using the command-line:

```
integrate gradient_file [-n nsteps] [-t temp] [-m (0|1)]  
                        [-h hill_height] [-f factor]
```

The gradient file name is provided first, followed by other parameters in any order. They are described below, with their default value in square brackets:

- -n: number of M-C steps to be performed; by default, a minimal number of steps is chosen based on the size of the grid, and the integration runs until a convergence criterion is satisfied (based on the RMSD between the target gradient and the real PMF gradient)
- -t: temperature for M-C sampling (unrelated to the simulation temperature) [500 K]
- -m: use metadynamics-like biased sampling? (0 = false) [1]
- -h: increment for the history-dependent bias (“hill height”) [0.01 kcal/mol]
- -f: if non-zero, this factor is used to scale the increment stepwise in the second half of the M-C sampling to refine the free energy estimate [0.5]

Using the default values of all parameters should give reasonable results in most cases.

abf_integrate produces the following output files:

- gradient_file.pmf: computed free energy surface
- gradient_file.histo: histogram of M-C sampling (not usable in a straightforward way if the history-dependent bias has been applied)
- gradient_file.est: estimated gradient of the calculated free energy surface (from finite differences)
- gradient_file.dev: deviation between the user-provided numerical gradient and the actual gradient of the calculated free energy surface. The RMS norm of this vector field is used as a convergence criteria and displayed periodically during the integration.

Note: Typically, the “deviation” vector field does not vanish as the integration converges. This happens because the numerical estimate of the gradient does not exactly derive from a potential, due to numerical approximations used to obtain it (finite sampling and discretization on a grid).

6.2 Metadynamics

The metadynamics method uses a history-dependent potential [9] that generalizes to any type of colvars the conformational flooding [10] and local elevation [11] methods, originally formulated to use as colvars the principal components of a covariance matrix or a set of dihedral angles, respectively. The metadynamics potential on the colvars $\xi = (\xi_1, \xi_2, \dots, \xi_{N_{cv}})$ is defined as:

$$V_{\text{meta}}(\xi) = \sum_{t'=\delta t, 2\delta t, \dots}^{t'<t} W \prod_{i=1}^{N_{cv}} \exp\left(-\frac{(\xi_i - \xi_i(t'))^2}{2\delta_{\xi_i}^2}\right), \quad (17)$$

where V_{meta} is the history-dependent potential acting on the *current* values of the colvars ξ , and depends only parametrically on the *previous* values of the colvars. V_{meta} is constructed as a sum of N_{cv} -dimensional repulsive Gaussian “hills”, whose height is a chosen energy constant W , and whose centers are the previously explored configurations $(\xi(\delta t), \xi(2\delta t), \dots)$. Each Gaussian functions has a width of approximately $2\delta_{\xi_i}$ along the direction of the i -th colvar.

During the simulation, the system evolves towards the nearest minimum of the “effective” potential of mean force $\tilde{A}(\xi)$, which is the sum of the “real” underlying potential of mean force $A(\xi)$ and the metadynamics potential $V_{\text{meta}}(\xi)$. Therefore, at any given time the probability of observing the configuration ξ^* is proportional to $\exp(-\tilde{A}(\xi^*)/\kappa_B T)$: this is also the probability that a new Gaussian “hill” is added at

that configuration. If the simulation is run for a sufficiently long time, each local minimum is canceled out by the sum of the Gaussian “hill” functions. At that stage the “effective” potential of mean force $\tilde{A}(\xi)$ is constant, and $-V_{\text{meta}}(\xi)$ is an accurate estimator of the “real” potential of mean force $A(\xi)$, save for an additive constant:

$$A(\xi) \simeq -V_{\text{meta}}(\xi) + K \quad (18)$$

Assuming that the set of collective variables includes all relevant degrees of freedom, the predicted error of the estimate is a simple function of the correlation times of the colvars τ_{ξ_i} , and of the user-defined parameters W , δ_{ξ_i} and δt [12]. In typical applications, a good rule of thumb can be to choose the ratio $W/\delta t$ much smaller than $\kappa_B T/\tau_{\xi}$, where τ_{ξ} is the longest among ξ ’s correlation times: δ_{ξ_i} then dictates the resolution of the calculated PMF.

To enable a metadynamics calculation, a metadynamics block must be defined in the colvars configuration file. Its only mandatory keyword is the colvars option listing all the variables involved: multidimensional PMFs are obtained by the same metadynamics instance applied to all the colvars.

The parameters W and δt are specified by the keywords `hillWeight` and `newHillFrequency`, respectively. The values of these options are optimal for colvars with correlation times τ_{ξ} in the range of a few thousand simulation steps, typical of many biomolecular simulations:

- `hillWeight` \langle Height of each hill (kcal/mol) \rangle
Context: metadynamics
Acceptable values: positive decimal
Default value: 0.01
Description: This option sets the height W of the hills that are added during this run. Lower values provide more accurate sampling at the price of longer simulation times to complete a PMF calculation.
- `newHillFrequency` \langle Frequency of hill creation \rangle
Context: metadynamics
Acceptable values: positive integer
Default value: 1000
Description: This option sets the number of integration steps after which a new hill is added to the metadynamics potential. Its value determines the parameter δt in eq. 17. Higher values provide more accurate sampling at the price of longer simulation times to complete a PMF calculation.

It is the user’s responsibility to either leave `hillWeight` and `newHillFrequency` at their default values, or to change them to match the specifics of each system. The parameter δ_{ξ_i} is instead defined as approximately half the width of the corresponding colvar ξ_i (see 3.1).

6.2.1 Output files

When interpolating grids are enabled (default behavior), the PMF is written every `colvarsRestartFrequency` steps to the file `output.pmf`. The following two options allow to control this behavior and to visually track statistical convergence:

- `writeFreeEnergyFile` \langle Periodically write the PMF for visualization \rangle
Context: metadynamics
Acceptable values: boolean
Default value: on
Description: When `useGrids` and this option are on, the PMF is written every `colvarsRestartFrequency` steps.

- `saveFreeEnergyFile` 〈Keep all the PMF files〉

Context: metadynamics

Acceptable values: boolean

Default value: off

Description: When `writeFreeEnergyFile` and this option are on, the step number is included in the file name. Activating this option can be useful to follow more closely the convergence of the simulation, by comparing PMFs separated by short times.

Note: when Gaussian hills are deposited near `lowerBoundary` or `upperBoundary` (see 3.1) and interpolating grids are used (default behavior), their truncation can give rise to accumulating errors. In these cases, as a measure of fault-tolerance all Gaussian hills near the boundaries are included in the output state file, and are recalculated analytically whenever the colvar falls outside the grid's boundaries. (Such measure protects the accuracy of the calculation, and can only be disabled by `hardLowerBoundary` or `hardUpperBoundary`.) To avoid gradual loss of performance and growth of the state file, either one of the following solutions is recommended:

- enabling the option `expandBoundaries`, so that the grid's boundaries are automatically recalculated whenever necessary; the resulting .pmf will have its abscissas expanded accordingly;
- setting `lowerWall` and `upperWall` well within the interval delimited by `lowerBoundary` and `upperBoundary`.

6.2.2 Performance tuning

The following options control the computational cost of metadynamics calculations, but do not affect results. Default values are chosen to minimize such cost with no loss of accuracy.

- `useGrids` 〈Interpolate the hills with grids〉

Context: metadynamics

Acceptable values: boolean

Default value: on

Description: This option discretizes all hills for improved performance, accumulating their energy and their gradients on two separate grids of equal spacing. Grids are defined by the values of `lowerBoundary`, `upperBoundary` and `width` for each colvar. Currently, this option is implemented for all types of variables except the non-scalar types (`distanceDir` or `orientation`). If `expandBoundaries` is defined in one of the colvars, grids are automatically expanded along the direction of that colvar.

- `hillWidth` 〈Relative width of the hills〉

Context: metadynamics

Acceptable values: positive decimal

Default value: $\sqrt{2\pi}/2$

Description: Along each colvar, the width of each Gaussian hill ($2\delta_{\xi_i}$) is given by the product between this number and the colvar's width. The default value gives hills whose volume is the product of W times the width of all colvars. For a smoother visualization of the free energy plot, decrease width and increase `hillWidth` in the same proportion. **Note:** when `useGrids` is on (default in most cases), values smaller than 1 should be avoided to avoid discretization errors.

- `rebinGrids` 〈Recompute the grids when reading a state file〉

Context: metadynamics

Acceptable values: boolean

Default value: off

Description: When restarting from a state file, the grid’s parameters (boundaries and widths) saved in the state file override those in the configuration file. Enabling this option forces the grids to match those in the current configuration file.

6.2.3 Well-tempered metadynamics

The following options define the configuration for the “well-tempered” metadynamics approach [13]:

- `wellTempered` `<Perform well-tempered metadynamics>`

Context: metadynamics

Acceptable values: boolean

Default value: off

Description: If enabled, this flag causes well-tempered metadynamics as described by Barducci et al. [13] to be performed, rather than standard metadynamics. The parameter `biasTemperature` is then required. This feature was contributed by Li Li (Luthey-Schulten group, Departement of Chemistry, UIUC).

- `biasTemperature` `<Temperature bias for well-tempered metadynamics>`

Context: metadynamics

Acceptable values: positive decimal

Description: When running metadynamics in the long time limit, collective variable space is sampled to a modified temperature $T + \Delta T$. In conventional metadynamics, the temperature “boost” ΔT would constantly increases with time. Instead, in well-tempered metadynamics ΔT must be defined by the user via `biasTemperature`. If `dumpFreeEnergyFile` is enabled, the written PMF includes the scaling factor $(T + \Delta T)/\Delta T$ [13]. A careful choice of ΔT determines the sampling and convergence rate, and is hence crucial to the success of a well-tempered metadynamics simulation.

6.2.4 Multiple-replicas metadynamics

The following options define metadynamics calculations with more than one replica:

- `multipleReplicas` `<Multiple replicas metadynamics>`

Context: metadynamics

Acceptable values: boolean

Default value: off

Description: If this option is on, multiple (independent) replica of the same system can be run at the same time, and their hills will be combined to obtain a single PMF [14]. Replicas are identified by the value of `replicaID`. Communication is done by files: each replica must be able to read the files created by the others, whose paths are communicated through the file `replicasRegistry`. This file, and the files listed in it, are read every `replicaUpdateFrequency` steps. Every time the colvars state file is written (`colvarsRestartFrequency`), the file:

“`output.colvars.name.replicaID.state`” is also written, containing the state of the metadynamics bias for `replicaID`. In the time steps between `colvarsRestartFrequency`, new hills are temporarily written to the file:

“`output.colvars.name.replicaID.hills`”, which serves as communication buffer. These files are only required for communication, and may be deleted after a new MD run is started with a different `outputName`.

- **replicaID** 〈Set the identifier for this replica〉
Context: metadynamics
Acceptable values: string
Description: If `multipleReplicas` is on, this option sets a unique identifier for this replica. All replicas should use identical collective variable configurations, except for the value of this option.
- **replicasRegistry** 〈Multiple replicas database file〉
Context: metadynamics
Acceptable values: UNIX filename
Default value: “*name.replica_files.txt*”
Description: If `multipleReplicas` is on, this option sets the path to the replicas’ database file.
- **replicaUpdateFrequency** 〈How often hills are communicated between replicas〉
Context: metadynamics
Acceptable values: positive integer
Default value: `newHillFrequency`
Description: If `multipleReplicas` is on, this option sets the number of steps after which each replica (re)reads the other replicas’ files. The lowest meaningful value of this number is `newHillFrequency`. If access to the file system is significantly affecting the simulation performance, this number can be increased, at the price of reduced synchronization between replicas. Values higher than `colvarsRestartFrequency` may not improve performance significantly.
- **dumpPartialFreeEnergyFile** 〈Periodically write the contribution to the PMF from this replica〉
Context: metadynamics
Acceptable values: boolean
Default value: on
Description: When `multipleReplicas` is on, the file *output.pmf* contains the combined PMF from all replicas. Enabling this option produces an additional file *output.partial.pmf*, which can be useful to quickly monitor the contribution of each replica to the PMF. The requirements for this option are the same as `dumpFreeEnergyFile`.

6.2.5 Compatibility and post-processing

The following options may be useful only for applications that go beyond the calculation of a PMF by metadynamics:

- **name** 〈Name of this metadynamics instance〉
Context: metadynamics
Acceptable values: string
Default value: “meta” + rank number
Description: This option sets the name for this metadynamics instance. While it is not advisable to use more than one metadynamics instance within the same simulation, this allows to distinguish each instance from the others. If there is more than one metadynamics instance, the name of this bias is included in the metadynamics output file names, such as e.g. the *.pmf* file.
- **keepHills** 〈Write each individual hill to the state file〉
Context: metadynamics
Acceptable values: boolean
Default value: off
Description: When `useGrids` and this option are on, all hills are saved to the state file in their

analytic form, alongside their grids. This makes it possible to later use exact analytic Gaussians for rebinGrids. To only keep track of the history of the added hills, writeHillsTrajectory is preferable.

- writeHillsTrajectory <Write a log of new hills>

Context: metadynamics

Acceptable values: boolean

Default value: on

Description: If this option is on, a logfile is written by the metadynamics bias, with the name “*output.colvars.<name>.hills.traj*”, which can be useful to follow the time series of the hills. When multipleReplicas is on, its name changes to “*output.colvars.<name>.<replicaID>.hills.traj*”. This file can be used to quickly visualize the positions of all added hills, in case newHillFrequency does not coincide with colvarsRestartFrequency.

6.3 Harmonic restraints

The harmonic biasing method may be used to enforce fixed or moving restraints, including variants of Steered and Targeted MD. Within energy minimization runs, it allows for restrained minimization, e.g. to calculate relaxed potential energy surfaces. In the context of the colvars module, harmonic potentials are meant according to their textbook definition: $V(x) = \frac{1}{2}k(x - x_0)^2$. Note that this differs from harmonic bond and angle potentials in common force fields, where the factor of one half is typically omitted, resulting in a non-standard definition of the force constant. A harmonic restraint is set up by a harmonic { ... } block, which may contain (in addition to the standard option colvars) the following keywords:

- forceConstant <Scaled force constant (kcal/mol)>

Context: harmonic

Acceptable values: positive decimal

Default value: 1.0

Description: This defines a scaled force constant for the harmonic potential. To ensure consistency for multidimensional restraints, it is divided internally by the square of the specific width for each colvar involved (which is 1 by default), so that all colvars are effectively dimensionless and of commensurate size. For instance, setting a scaled force constant of 10 kcal/mol acting on two colvars, an angle with a width of 5 degrees and a distance with a width of 0.5 Å, will apply actual force constants of $0.4 \text{ kcal/mol} \times \text{degree}^{-2}$ for the angle and $40 \text{ kcal/mol}/\text{\AA}^2$ for the distance.

- centers <Initial harmonic restraint centers>

Context: harmonic

Acceptable values: space-separated list of colvar values

Description: The centers (equilibrium values) of the restraint are entered here. The number of values must be the number of requested colvars. Each value is a decimal number if the corresponding colvar returns a scalar, a “(x, y, z)” triplet if it returns a unit vector or a vector, and a “q0, q1, q2, q3)” quadruplet if it returns a rotational quaternion. If a colvar has periodicities or symmetries, its closest image to the restraint center is considered when calculating the harmonic potential.

Tip: A complex set of restraints can be applied to a system, by defining several colvars, and applying one or more harmonic restraints to different groups of colvars. In some cases, dozens of colvars can be defined, but their value may not be relevant: to limit the size of the colvars trajectory file, it may be wise to disable outputValue for such “ancillary” variables, and leave it enabled only for “relevant” ones.

6.3.1 Moving restraints: steered molecular dynamics

The following options allow to change gradually the centers of the harmonic restraints during a simulations. When the centers are changed continuously, a steered MD in a collective variable space is carried out.

- **targetCenters** 〈Steer the restraint centers towards these targets〉
Context: harmonic
Acceptable values: space-separated list of colvar values
Description: When defined, the current centers will be moved towards these values during the simulation. By default, the centers are moved over a total of `targetNumSteps` steps by a linear interpolation, in the spirit of Steered MD. If `targetNumStages` is set to a nonzero value, the change is performed in discrete stages, lasting `targetNumSteps` steps *each*. This second mode may be used to sample successive windows in the context of an Umbrella Sampling simulation. When continuing a simulation run, the centers specified in the configuration file `<colvarsConfig>` are overridden by those saved in the restart file `<colvarsInput>`. To perform Steered MD in an arbitrary space of colvars, it is sufficient to use this option and enable `outputAppliedForce` within each of the colvars involved.
- **targetNumSteps** 〈Number of steps for steering〉
Context: harmonic
Acceptable values: positive integer
Description: In single-stage (continuous) transformations, defines the number of MD steps required to move the restraint centers (or force constant) towards the values specified with `targetCenters` or `targetForceConstant`. After the target values have been reached, the centers (resp. force constant) are kept fixed. In multi-stage transformations, this sets the number of MD steps *per stage*.
- **outputCenters** 〈Write the current centers to the trajectory file〉
Context: harmonic
Acceptable values: boolean
Default value: off
Description: If this option is chosen and `colvarsTrajFrequency` is not zero, the positions of the restraint centers will be written to the trajectory file during the simulation. This option allows to conveniently extract the PMF from the colvars trajectory files in a steered MD calculation.
- **outputAccumulatedWork** 〈Write the accumulated work to the trajectory file〉
Context: harmonic
Acceptable values: boolean
Default value: off
Description: If this option is chosen and `colvarsTrajFrequency` is not zero, the accumulated work from the beginning of the simulation will be written to the trajectory file. If the simulation has been continued from a previous state file, the previously accumulated work is included in the integral. This option allows to conveniently extract the PMF from the colvars trajectory files in a steered MD calculation.

Note on restarting moving restraint simulations: Information about the current step and stage of a simulation with moving restraints is stored in the restart file (state file). Thus, such simulations can be run in several chunks, and restarted directly using the same colvars configuration file. In case of a restart, the values of parameters such as `targetCenters`, `targetNumSteps`, etc. should not be changed manually.

6.3.2 Moving restraints: umbrella sampling

The centers of the harmonic restraints can also be changed in discrete stages: in this cases a one-dimensional umbrella sampling simulation is performed. The sampling windows in simulation are calculated in sequence. The colvars trajectory file may then be used both to evaluate the correlation times between consecutive windows, and to calculate the frequency distribution of the colvar of interest in each window. Furthermore, frequency distributions on a predefined grid can be automatically obtained by using the histogram bias (see 6.4).

To activate an umbrella sampling simulation, the same keywords as in the previous section can be used, with the addition of the following:

- **targetNumStages** \langle Number of stages for steering \rangle
Context: harmonic
Acceptable values: non-negative integer
Default value: 0
Description: If non-zero, sets the number of stages in which the restraint centers or force constant are changed to their target values. If zero, the change is continuous. Each stage lasts targetNumSteps MD steps. To sample both ends of the transformation, the simulation should be run for targetNumSteps \times (targetNumStages + 1).

6.3.3 Changing force constant

The force constant of the harmonic restraint may also be changed to equilibrate [?].

- **targetForceConstant** \langle Change the force constant towards this value \rangle
Context: harmonic
Acceptable values: positive decimal
Description: When defined, the current forceConstant will be moved towards this value during the simulation. Time evolution of the force constant is dictated by the targetForceExponent parameter (see below). By default, the force constant is changed smoothly over a total of targetNumSteps steps. This is useful to introduce or remove restraints in a progressive manner. If targetNumStages is set to a nonzero value, the change is performed in discrete stages, lasting targetNumSteps steps *each*. This second mode may be used to compute the conformational free energy change associated with the restraint, within the FEP or TI formalisms. For convenience, the code provides an estimate of the free energy derivative for use in TI. A more complete free energy calculation (particularly with regard to convergence analysis), while not handled by the colvars module, can be performed by post-processing the colvars trajectory, if colvarsTrajFrequency is set to a suitably small value. It should be noted, however, that restraint free energy calculations may be handled more efficiently by an indirect route, through the determination of a PMF for the restrained coordinate.[?]
- **targetForceExponent** \langle decimal equal to or greater than 1.0 \rangle
Context: Exponent in the time-dependence of the force constant
Acceptable values: 1.0
Default value: Sets the exponent, α , in the function used to vary the force constant as a function of time. The force is varied according to a coupling parameter λ , raised to the power α : $k_\lambda = k_0 + \lambda^\alpha(k_1 - k_0)$, where k_0 , k_λ , and k_1 are the initial, current, and final values of the force constant. The parameter λ evolves linearly from 0 to 1, either smoothly, or in targetNumStages equally spaced discrete stages, or according to an arbitrary schedule set with lambdaSchedule. When the initial value of the force constant is zero, an exponent greater than 1.0 distributes the effects of introducing the

restraint more smoothly over time than a linear dependence, and ensures that there is no singularity in the derivative of the restraint free energy with respect to λ . A value of 4 has been found to give good results in some tests.

Description:

- `targetEquilSteps` \langle Number of steps discarded from TI estimate \rangle
Context: harmonic
Acceptable values: positive integer
Description: Defines the number of steps within each stage that are considered equilibration and discarded from the restraint free energy derivative estimate reported in the output.
- `lambdaSchedule` \langle Schedule of λ -points for changing force constant \rangle
Context: harmonic
Acceptable values: list of real numbers between 0 and 1
Description: If specified together with `targetForceConstant`, sets the sequence of discrete λ values that will be used for different stages.

6.4 Multidimensional histograms

The histogram feature is used to record the distribution of a set of collective variables in the form of a N-dimensional histogram. It functions as a “collective variable bias”, and is invoked by adding a histogram block to the *colvars* configuration file.

In addition to the common parameters `name` and `colvars` described above, a histogram block may define the following parameter:

- `outputFreq` \langle Frequency (in timesteps) at which the histogram file is refreshed \rangle
Context: histogram
Acceptable values: positive integer
Default value: Colvar module restart frequency
Description: The file containing histogram data is written on disk at the given time interval.

Like the ABF and metadynamics biases, histogram uses parameters from the colvars to define its grid. The grid ranges from `lowerBoundary` to `upperBoundary`, and the bin width is set by the `width` parameter.

References

- [1] M. Iannuzzi, A. Laio, and M. Parrinello. Efficient exploration of reactive potential energy surfaces using car-parrinello molecular dynamics. *Phys. Rev. Lett.*, 90(23):238302, 2003. 11, 21
- [2] E A Coutsiias, C Seok, and K A Dill. Using quaternions to calculate RMSD. *J. Comput. Chem.*, 25(15):1849–1857, 2004. 15, 22, 25
- [3] Eric Darve, David Rodríguez-Gómez, and Andrew Pohorille. Adaptive biasing force method for scalar and vector free energy calculations. *J. Chem. Phys.*, 128(14):144120, 2008. 29
- [4] J. Hénin and C. Chipot. Overcoming free energy barriers using unconstrained molecular dynamics simulations. *J. Chem. Phys.*, 121:2904–2914, 2004. 29
- [5] A. Carter, E. G. Ciccotti, J. T. Hynes, and R. Kapral. Constrained reaction coordinate dynamics for the simulation of rare events. *Chem. Phys. Lett.*, 156:472–477, 1989. 29
- [6] M. J. Ruiz-Montero, D. Frenkel, and J. J. Brey. Efficient schemes to compute diffusive barrier crossing rates. *Mol. Phys.*, 90:925–941, 1997. 29
- [7] W. K. den Otter. Thermodynamic integration of the free energy along a reaction coordinate in cartesian coordinates. *J. Chem. Phys.*, 112:7283–7292, 2000. 29
- [8] Giovanni Ciccotti, Raymond Kapral, and Eric Vanden-Eijnden. Blue moon sampling, vectorial reaction coordinates, and unbiased constrained dynamics. *ChemPhysChem*, 6(9):1809–1814, 2005. 29
- [9] A. Laio and M. Parrinello. Escaping free-energy minima. *Proc. Natl. Acad. Sci. USA*, 99(20):12562–12566, 2002. 33
- [10] Helmut Grubmüller. Predicting slow structural transitions in macromolecular systems: Conformational flooding. *Phys. Rev. E*, 52(3):2893–2906, Sep 1995. 33
- [11] T. Huber, A. E. Torda, and W.F. van Gunsteren. Local elevation - A method for improving the searching properties of molecular-dynamics simulation. *Journal of Computer-Aided Molecular Design*, 8(6):695–708, DEC 1994. 33
- [12] G. Bussi, A. Laio, and M. Parrinello. Equilibrium free energies from nonequilibrium metadynamics. *Phys. Rev. Lett.*, 96(9):090601, 2006. 34
- [13] Alessandro Barducci, Giovanni Bussi, and Michele Parrinello. Well-tempered metadynamics: A smoothly converging and tunable free-energy method. *Phys. Rev. Lett.*, 100:020603, 2008. 36
- [14] P. Raiteri, A. Laio, F. L. Gervasio, C. Micheletti, and M. Parrinello. Efficient reconstruction of complex free energy landscapes by multiple walkers metadynamics. *J. Phys. Chem. B*, 110(8):3533–9, 2005. 36